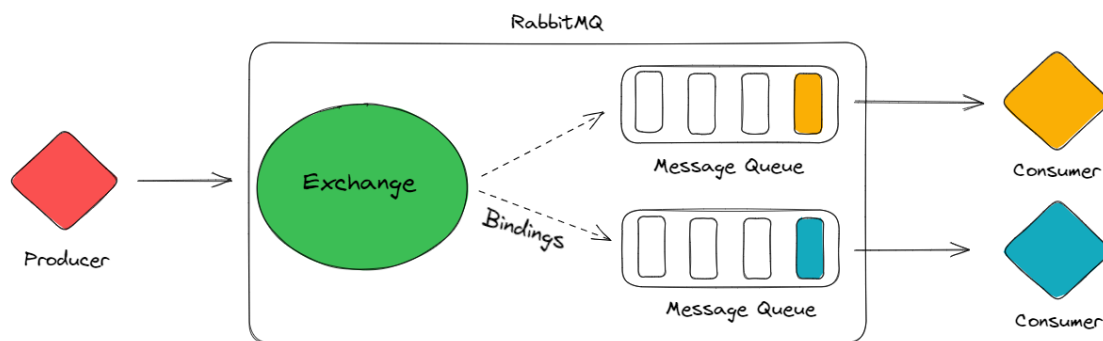


# Microservices & RabbitMQ lab Report

Amo Samuel

## Brief Overview of RabbitMQ

RabbitMQ is a widely-used messaging broker that allows applications to communicate asynchronously and decouples processes through message queues. Its architecture is built around the following core concepts.



## Key Components:

1. **Producer:** The application that sends the message. It does not communicate directly with the consumer but instead with an **Exchange**.
2. **Exchange:** This component receives the message from the producer and determines which queue(s) to route the message to, based on routing rules and bindings.
3. **Bindings:** These are the rules that tell the exchange which queue(s) to send the message to.
4. **Message Queue:** A place where messages are stored until they are consumed by the consumer. The queue is independent of both the producer and the consumer, providing reliability and asynchronous messaging.
5. **Consumer:** The application that listens for messages from the queue, processes them, and can save or perform any logic with the message.

RabbitMQ acts as the intermediary, allowing these components to interact while remaining decoupled, which increases the scalability and flexibility of the system.

## My RabbitMQ Configuration in Code

In my project, I configured a Spring Boot application to use RabbitMQ for order processing. Below is how I went about implementing RabbitMQ's messaging architecture using Spring AMQP.

### 1. Defining the Configuration:

In the `RabbitMQConfigurations` class, I used Spring's `@Configuration` to define essential RabbitMQ components such as the queue, exchange, and bindings. The configuration details were fetched from application properties using `@Value`.

```
@Configuration
public class RabbitMQConfigurations {

    @Value("${spring.rabbitmq.queue}")
    public String orderQueue;
    @Value("${spring.rabbitmq.exchange}")
    public String orderExchange;
    @Value("${spring.rabbitmq.routing-key}")
    public String orderRoutingKey;
```

- **Queue:** I declared the order queue using the `Queue` class.
- **Exchange:** The `TopicExchange` was set up to route the messages based on a routing key.
- **Binding:** This connects the queue to the exchange using the routing key provided.

The binding ensures that any messages sent to the exchange with the correct routing key are delivered to the intended queue

## 2. Message Conversion and Sending:

To send messages in JSON format, I used the `Jackson2JsonMessageConverter` to convert the message before sending it to the RabbitMQ exchange.

The `RabbitTemplate` was configured to use this message converter to simplify message processing between the producer and consumer.

```
@Bean
public RabbitTemplate rabbitTemplate(ConnectionFactory connectionFactory) {
    RabbitTemplate rabbitTemplate = new RabbitTemplate(connectionFactory);
    rabbitTemplate.setMessageConverter(jsonMessageConverter());
    return rabbitTemplate;
}

@Bean
public Jackson2JsonMessageConverter jsonMessageConverter() {
    return new Jackson2JsonMessageConverter();
}
```

## 3. OrderProducer - Sending Messages:

The `OrderProducer` is responsible for sending messages (orders) to RabbitMQ. I used the `RabbitTemplate` to send the message to the exchange with the appropriate routing key.

```
@PostMapping
public String sendOrder(@RequestBody Order order) {
    if (order == null || order.getOrderId() == null) {
        throw new OrderProcessingException("Invalid Order Data");
    }
    rabbitTemplate.convertAndSend(rabbitMQConfigurations.orderExchange, rabbitMQConfigurations.orderRoutingKey, order);
    return "Order sent successfully";
}
```

The order is validated before sending, and any issues (such as missing order IDs) trigger a custom exception (`OrderProcessingException`), ensuring the system handles invalid data.

## OrderConsumer - Receiving Messages:

The `OrderConsumer` class listens to the queue for incoming messages using the `@RabbitListener` annotation. Once a message is received, it is logged and saved to the database using the `OrderRepository`.

```
    }  
  
    @RabbitListener(queues = "${spring.rabbitmq.queue}")  
    @ public void processOrder(Order order) {  
        logger.info("Order received and saving to database {}", order.toString());  
  
        orderRepository.save(order);  
    }  
}
```

The consumer logs the received order and processes it by saving the data into the database. This method demonstrates the asynchronous nature of the message broker, as it allows orders to be processed independently of the producer.

## 5. Error Handling:

To ensure that the application remains stable in case of errors, I implemented a global exception handler. The `GlobalExceptionHandler` captures any exceptions that occur while processing an order and provides a detailed error response to the client.

```
private final Logger logger = LoggerFactory.getLogger(GlobalExceptionHandler.class); 2 usages  
  
@ExceptionHandler(OrderProcessingException.class)  
public ResponseEntity<String> handleOrderProcessingException(OrderProcessingException processingException) {  
    logger.error("Error Processing Order {}", processingException.getMessage());  
    return ResponseEntity.status(HttpStatus.INTERNAL_SERVER_ERROR).body(processingException.getMessage());  
}
```

With this setup, the application gracefully handles order processing exceptions, logging errors, and returning informative error messages without crashing.

## 6. Data Persistence:

The `Order` entity is annotated with JPA annotations, and the `OrderRepository` is responsible for saving orders to the database. This separation of concerns ensures that the producer and consumer deal only with messaging, while the repository handles data persistence.

```
@Entity
@Entity
@Table(name = "orders")
public class Order {
    @Id
    private String orderId;
    private String productName;
    private int quantity;
}
```

## Conclusion:

Using RabbitMQ, I was able to decouple the producer and consumer of orders, allowing for scalable and fault-tolerant order processing. The architecture ensures that messages can be processed asynchronously, with a strong focus on reliable error handling and logging. The system is now capable of efficiently processing orders, regardless of the speed at which they are produced or consumed.